Aviation Accident Analysis: Trends, Causes, and Safety Measures from 1962 to 2022

Taran Mamidala (50604177)
Nanda Krishna Sindam (50604161)
Sai Dinesh Rachakonda (50533608)

***Abstract*—In this project, we have taken a dataset on airplane crashes occurred worldwide from 1948 to 2022. The dataset is taken from Kaggle website. We have performed the 10 distinct Data Cleaning steps and 10 distinct Exploratory Data Analysis (EDA) steps on the dataset.**

Keywords—Airplane accidents, Safety

I. *Background and Problem Statement*

Aviation accidents, though rare, significantly affect safety, security, and the entire transportation sector as a whole. Understandably, the nature or cause of such accidents is very important to know in order to avoid similar incidents in the near future so that the passengers' and crew's safety may be warranted. The aviation sector is highly governed by laws and regulations, and even small mistakes can cost one's life, financial loss to airlines, legal issues, and erosion of public trust.

Even with tremendous improvement in technology, safety, and training within the aviation sector over the years, accidents do occur occasionally. These are caused by many complex factors, which include mechanical failure, human error, bad weather, and environmental problems. Each failure provides an important opportunity to review information, identify trends, and improve safety procedures.

Objectives:

Analysis of Data and Insights: To Identify trends in aviation accidents in order to make useful inferences from the given dataset regarding causal factors associated with weather conditions, human error, technical failure, and other contributing factors.

Visualization of Locations of Accidents: To Avail visualizations to project accident locations onto a world map that may assist in detecting spatial patterns in accidents. This could delineate areas around the world where the incidence of aviation accidents is higher compared to other areas, thus suggesting areas to which targeted safety interventions might most usefully be directed.

Identification of Probable Causes: Investigate the probable causes of these accidents and classify them so that better understanding of the common factors can be identified. These will also enable stakeholders to give emphasis on those areas that need priority improvement.

Contribution to Safety Protocols

Contribute to aviation authorities, airlines, and safety organizations by providing inputs on enhancing safety protocols in the form of precautionary measures and enhancement of pilot and crew training. Problem Significance Addressing aviation accidents comes with saving lives and improving air travel safety. A single aviation accident can have disastrous consequences in human lives and finances. Studying the accidents that have occurred over the years can help stakeholders discover the trends to build mitigation strategies that will help improve safety and meet regulations.

In addition, the findings from this analysis may help inform public policy and guarantee technological innovation in aviation safety, further solidifying the industry of aviation. Air travel is a critical aspect in developing global relationships and business; therefore, ensuring its safety involves not only meeting regulatory requirements but also a moral obligation to ensure safety for both passengers and crews.

In other words, the aviation accident data set provides an overview and a method to try and understand the complexity of aviation safety; thus, it becomes one of the major issues today.

A. *The Potential of this Dataset*

The study of this aviation disaster data set may greatly improve the knowledge of what factors lead to aviation accidents and what can be done to ensure greater safety in the field of aviation. So, here's how this project can add to the problem field:

Enhanced Data-Driven Decision Making: 1 Because with all of this data, the airlines, FAA, and other safety organizations can make decisions based on facts and not just speculation. The insights gained from this analysis can lead to targeted safety interventions and policy adjustments that are grounded in data.

Identification of Risk Patterns: This project is designed to find common elements and danger factors of aviation accidents. By categorizing incidents based on factors like weather conditions, geographical locations, and types of aircraft involved, it can help aviation professionals understand where vulnerabilities lie. That information could lead to better training programs and safety precautions.

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

Development of Predictive Models: From the data, one could infer the probability of an accident occurring under certain conditions, and that would provide a basis for predictive models. These models will enable the airlines to forecast potential dangers and take the necessary precautions to ensure safety.

Improved Training and Protocols: The study of human error and procedural errors would allow us to create better training programs for pilots, crew, and ground personnel. That way all the staff would know what to do in the case of an emergency, and accidents would be less likely to happen.

Support for Regulatory Changes: Any results from this project can serve as grounds to lobby for changes in the laws and safety requirements. With the project providing proof of certain factors that may have caused accidents, the hope is that the policymakers will see this and make the regulations concerning these factors stricter in order to eliminate these risks.

Public Awareness and Trust: Dissemination of the results of the project can also serve to inform the public of the state of aviation safety and the strides that are continually being taken to ensure it. That openness could foster a lot of trust in the passengers, letting them know that the airline industry really does care about their safety.

Why This Contribution is Crucial

Human Lives at Stake: A plane crash is no joke, loss of life, and shattering the lives of families and communities. If the project can in some way help to improve safety, then who knows, lives may be saved and tragedies averted.

Financial Implications: Not to mention accidents would cost an airline millions of dollars in lawsuits, insurance claims, and loss of company image. Well, safety precautions could be increased to minimize those risks, and the financial state of the aviation industry would become more stable.

Global Connectivity: The airline industry is a backbone of international accessibility allowing travel and business around the world. Safety is not only good for the industry itself, but for global relations and economic development.

Continuous Improvement in Aviation Safety: The field of aviation has come a long way in safety in the last few decades, but there is still so much that can be done. It adds to the discussion of the safety of aviation and allows the industry to change and grow in response to these new issues.

Informed Policy and Practice: The insights derived from this project can lead to evidence-based policies and practices, ensuring that safety measures are not just reactive but proactive. This helps to create a safety culture that exists at every level in aviation.

All in all, this project could really contribute to the safety of aviation by providing feedback that could lead to better training, regulations, and general safety precautions. The importance of these contributions cannot be understated considering that aviation safety affects human lives, economic stability, and the connectiveness of the world.

Data Sources

We used the NTSB Aviation Accident Dataset which contains detailed records of accidents from 1962 to 2023. This dataset includes over 80,000 records and contains information such as event dates, locations, accident types, aircraft damage, total fatalities, and environmental factors like weather conditions. The dataset is large and detailed enough to perform significant statistical analysis and predictive modeling. The link for the dataset is: https://www.kaggle.com/datasets/khsamaha/aviation-accident-database-synopses?select=AviationData.csv[1].

Dataset Characteristics:

- Size: 88889 Records, 31 Columns
- Important Attributes: Investigation Type, Event Date, Location, Country, Latitude, Longitude, Make, Model, Engine Type, Injury Severity, Aircraft Damage, Total Fatal Injuries, Weather Condition, etc.

Data Cleaning

1.Handling Missing Values:
What was done: We removed rows where missing values occurred jointly in columns named Location, Latitude, and Longitude as the rows with missing values are not necessary.

Why Necessary: Missing values can be used to introduce errors in any analysis or result in incorrect results. By removing the missing values with excessive missing data, we ensure the dataset is complete and reliable because in performing the spatial analysis of accidents. Those rows cannot be used if both coordinates and locations are missing

2. Handling Rows with Missing Values in Location, Latitude, and Longitude:
What We Did: For all rows where Location was missing but coordinates were available, we filled the Location with the name of the city for those coordinates.
Why It's Needed: Finding the available coordinates to fill in the missing names of locations will enable us to preserve much valuable data that would otherwise have to be discarded and allow for a more comprehensive variety of analyses. This ensures that location data is full and geographic mapping of aviation accidents can be done accurately.
This extra step in filling the missing values in Location due to coordinates actually increases the quality of the data because more granular insights can be achieved, especially in visualizations and spatial trend analysis.

3. Duplicate Removal:
What We Did: We checked for duplicate records and removed them so that the dataset had unique entries.
Why Important: The repetition biases the analysis towards an incorrect over-representation of the frequency of occurrence.

Think of a specific accident repeating several times in the dataset. This will skew any statistical model, even those based on machine learning, that would later be developed. Duplicate removal cleans the data.

4.Removing Numerical values from Injury severity column:

What We Did: We cleaned the "Injury Severity" column by removing any numerical values that had been entered into this column by mistake; this column should only contain valid categorical values, such as "Fatal" or "Non-Fatal."
Why it's Necessary: "Injury Severity" is a qualitative measure of data; therefore, there should not be any numerical values inside this column. Removing these erroneous values allows the data to remain consistent and, hence, the severity classification more valid and accurate for analysis.

5. Impute "Non-Fatal" in Injury Severity Column Where Total Fatal Injuries are 0:
We did the following: We replaced blank values in the column "Injury Severity" with "Non-Fatal" for those rows where "Total Fatal Injuries" was 0 but the "Injury Severity" column was blank.
Why it is necessary: This step makes the data uniform and logically consistent. In this context, blank values may mislead any model or analyst, but their filling on the basis of actual injury data enhances the dataset's completeness and accuracy.

6. Convert Data Types:
What We Did: We changed the data type of columns like event date to datetime and ensured that numeric columns were correctly cast to integers or floats.
Why it's Necessary: The use of proper datatypes is highly important in analysis. For example, there may come a situation where an incorrect datatype might make the code fail during filtering of dates or calculation over numerical values. The conversion to the right types further makes the analysis of data comfortable.

7. Standards for Column Naming:
What We Did: We standardized the column naming by making all cases lower and replacing spaces with underscores.
Why it is Necessary: Standardized column names make working programmatically with the dataset much easier, writing code for transformations, merging, or visualizations. Consistent naming reduces mistakes and improves readability of code.

8. Feature Engineering:
What We Did: We created the event_year, event_month, and event_day features from the event_date column for time-based analyses.
Why It's Needed: Such separation of event_date into various year, month, and day components enables more granular time-based analysis. This is useful when analyzing trends-for example, whether or not accidents are more frequent in specific years, months, or even days of the week.

9. Cleaning up the text columns by removing unwanted characters:
What We Did: Data cleaning on text columns was done by erasing unwanted characters such as special symbols or extra spaces.
Why It's Necessary: Text cleaning allows data to remain consistent. The data will not parse errors, allowing for better text-based analysis and filtering. Clean text really gives great quality to any text analysis or natural language processing that may be conducted on the dataset.

10.Binning Numerical Data

What we did: Since the range between minimum and maximum values of total_fatal_injuries is huge, we have binned these data into discrete bins for a better study of data. So, we have created bins like 0, 1-5, 6-10, 11-20, and 20+. We also added one binary column named HasFatalInjury which tells us if an incident has any fatality or not.

Why we did: The binning process here really allowed us to reduce the numerical data in this case into categorical representations that were far easier to analyze and visually present the injury severity across the accidents. This also brought out some trends and patterns in how injuries are related with other variables.

Usage: The variable injury_category as categorized herein assists in the understanding of the frequency of accidents within specified ranges of injury, thus enabling necessary comparisons across different aircraft types or accident conditions. HasFatalInjury as an indicator is certainly useful in terms of predictive modeling; this allows for a concentration on incidents that are with fatalities in such a way that relevance and effectiveness might be enhanced in regard to our safety assessments and interventions.

11.OutlierDetection:
What we did: We performed an outlier detection on numeric columns such as Total.Fatal.Injuries and handled those outliers either by capping them or removing them depending on the context.
Why is it important? The outliers can bias statistical analyses and models and hence often lead to wrong results. For example, an exceptionally large number of fatalities caused by one accident may mislead the average injury rate for all accidents. Handling such outliers makes sure that the data is representative of typical behavior and suitable for analysis.

How to do: Last Step - Reset the Index
What did we do? At the end of these processes, we cleaned up the index of the DataFrame to keep this in an orderly structure. Why is this necessary? Resetting the index ensures that after cleaning, the DataFrame is well-structured without gaps or misalignment. This is because further analysis and visualization require a clean and organized dataset for proper processing.

12. Outlier Removal

What we did: After detecting outliers in various numerical columns (such as Total Fatal Injuries, Total Serious Injuries, and Total Undetected), we performed an outlier removal step. This involved dropping identified outlier rows to clean the dataset and ensure that it contained only valid and relevant records. Finally, we reset the index of the DataFrame.

Why we did: By removing outlier entries, we ensured that our dataset more accurately reflected typical behavior, reducing the potential for skewed results in subsequent analyses and modeling. Resetting the index after this operation allowed us to maintain a clean and orderly structure in our DataFrame, which is essential for further data processing and exploration.

Usage: This step is crucial for ensuring the integrity of the data for future analyses. It helps avoid biases that could arise from extreme values, thus enhancing the reliability of the conclusions drawn from the dataset. A clean index also simplifies referencing specific rows in the DataFrame, making it easier to manipulate the data for further exploratory data analysis (EDA) or modeling tasks.

Exploratory Data Analysis

Overview:

Exploratory Data Analysis is a very important activity in the field of data analysis. The task aims at summarizing the general characteristics of the dataset and is often performed visually. Following NIST and John Tukey principles, we have performed various EDA tasks on aviation accident data to get meaningful insights that guide subsequent modeling and analytics.

EDA Steps:

1. Analysis of Injury Severity

Operation: The incidents are categorized as Minor, Non-Fatal, and Serious according to their injury severity.

Outcome: It helped us to identify the most common injury types because of crashes.

Use: This information will be useful in future analyses, as it will help in focusing safety measures on certain injury types that can then drive feature selection for predictive modeling on injury severity.

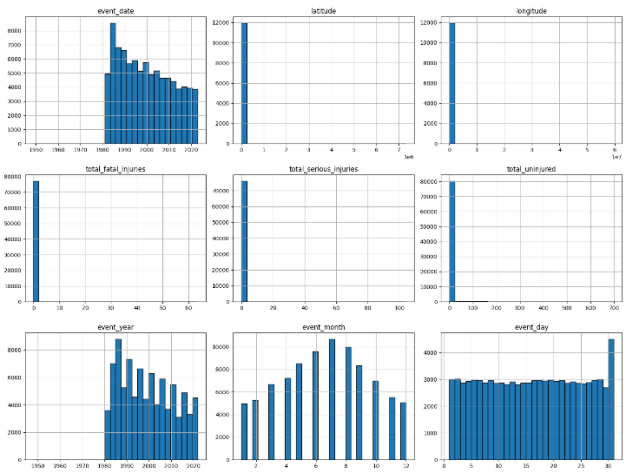

2. Annual Crash Trends

Operation: The total crashes per year and the top 15 airplane brands based on accident counts were mapped to bar graphs.

Outcome: Trends and spikes in the timings of accidents were detected, as did brands that have higher crash frequencies.

Usage: This time-series trend analysis enables stakeholders to be informed of times of year when flying can be more or less dangerous, which may spur some feature engineering based on time variables such as year and seasonality.

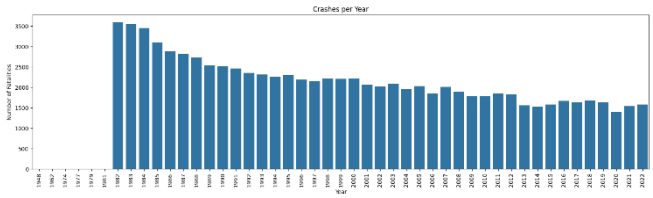

3. Fatal Injuries by Investigation Type

Operation: This created count plots to understand the number of different varieties of investigations.

Output: The types of investigations conducted told a lot about which kind of investigations were most common to identify areas of concern or regulatory foci.

Use: Investigation type knowledge assists in selecting appropriate features while undertaking predictive modeling for identifying factors that cause crashes.

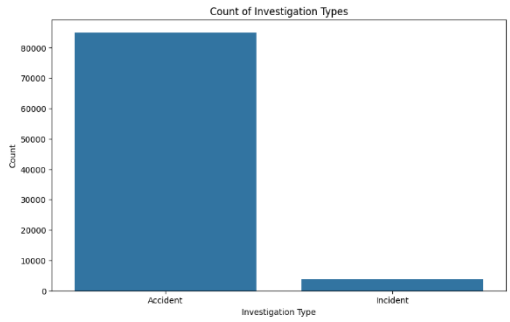

4. Aircraft Damage vs Injury Severity

Operation: Bar plots on Aircraft Damage vs Injury Severity

Output: A distribution of number of accidents in each category of aircraft damage as opposed to the severity of the injury was produced, hence showing the extent of the accidents' impact.

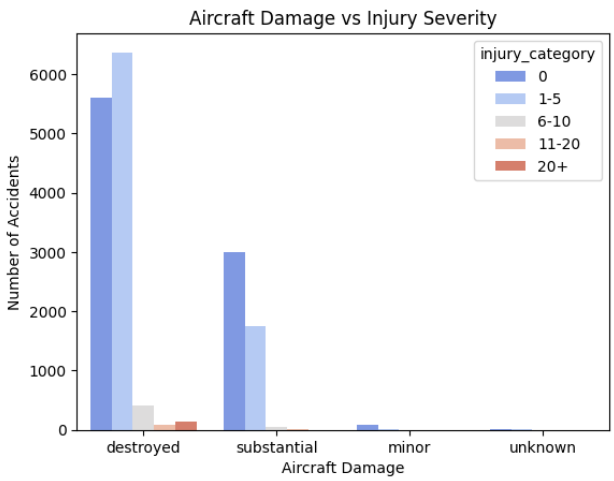

5. Crash Phases Analysis

Action: Created a bar plot of crashes based on the phase of flight-landing, takeoff, etc.

Result: Identified critical flight phases that have the highest accident rates.

Use: The idea is to inform feature engineering regarding flight phase and perhaps influence model inputs germane to specific operational risks.

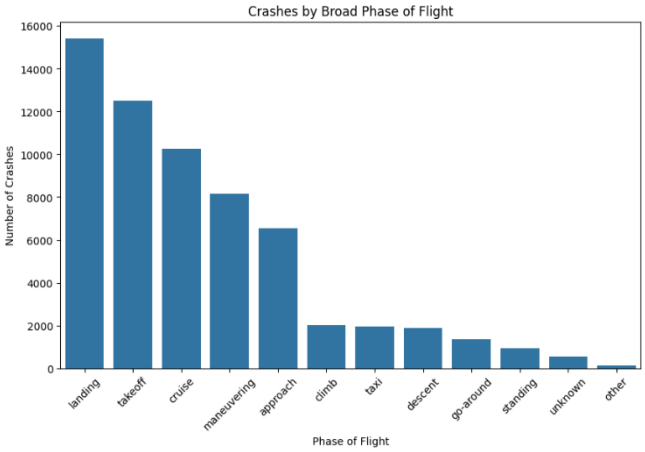

6. Geographic Distribution of Crashes

Operation: Plotted accident locations on a world map by latitude and longitude.

Output: Distributions of accidents, visualized over the world, identifying some geographical patterns.

Usage: Geographic insights can add the spatial feature to the work and help target where implementation of safety improvements may be needed in high-risk regions.

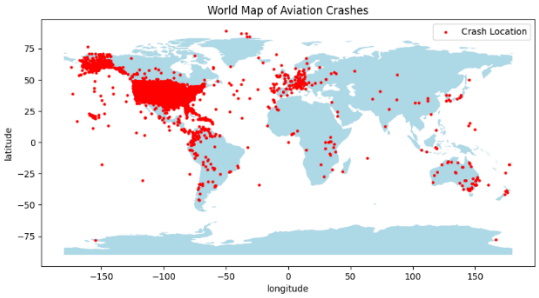

7. Weather Conditions Analysis

Operation: Created bar plots to analyze crashes against the weather conditions.

Outcome: The impact of weather on crash incidence was shown by finding the most frequent weather conditions that accidents are associated with.

Usage: Informs this analysis on exogenous variables that can be incorporated into predictive models since it established how significant environmental features can be.

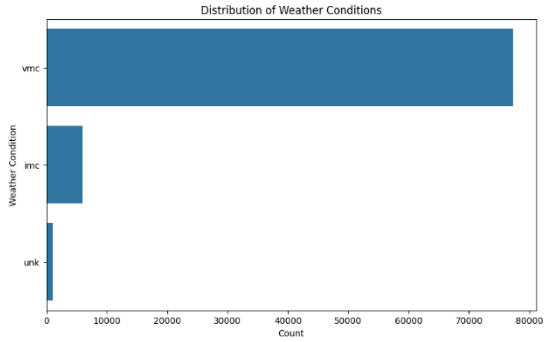

8. Line Plot for Total Fatal Injuries

Operation: Plotted line graphs showing total injuries by year.

Outcome: Extracted trend lines on variation of injury severity over years, which helped determine the change on either improvement or deterioration in safety.

Usage: This temporal analysis will inform the creation of time-series models that forecast future injury trends.

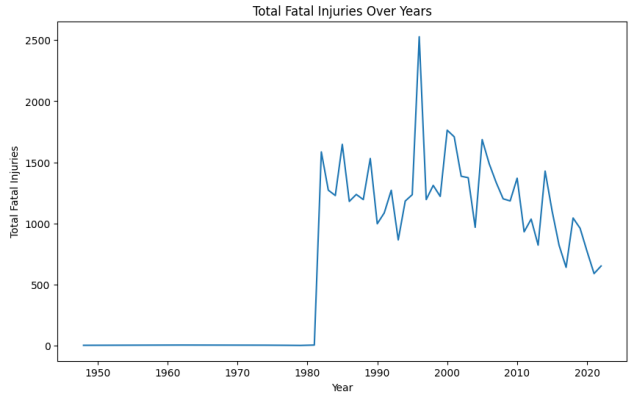

9. Aircraft Brand Analysis

Operation: Created bar plots for the top 10 aircraft brands by crash rate.

Outcome: Directed attention to those brands that require further safety assessments.

Usage: The output of this analysis will facilitate feature selection and model building aimed at brand-specific safety assessments.

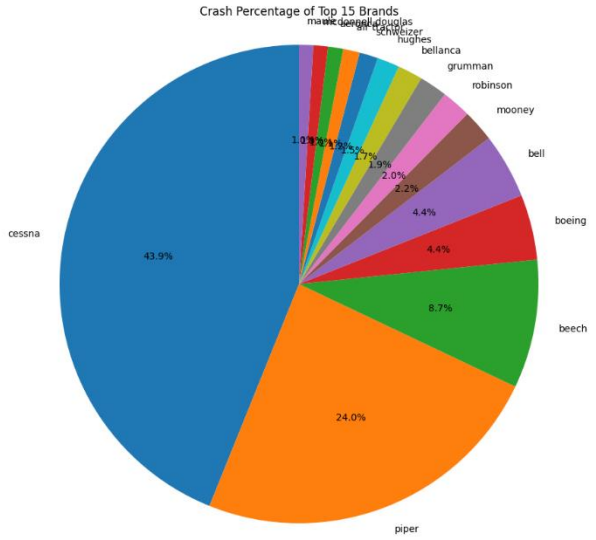

10. Correlation Matrix

Operation: Performed a correlation matrix analysis to examine the relationships between numerical features.

Outcome: Emphasized multicollinearity and feature interaction, which is relevant and may influence model performance.

Usage: The correlation insights are very helpful in selecting those features that are meaningfully contributing to a predictive model while avoiding redundancy.

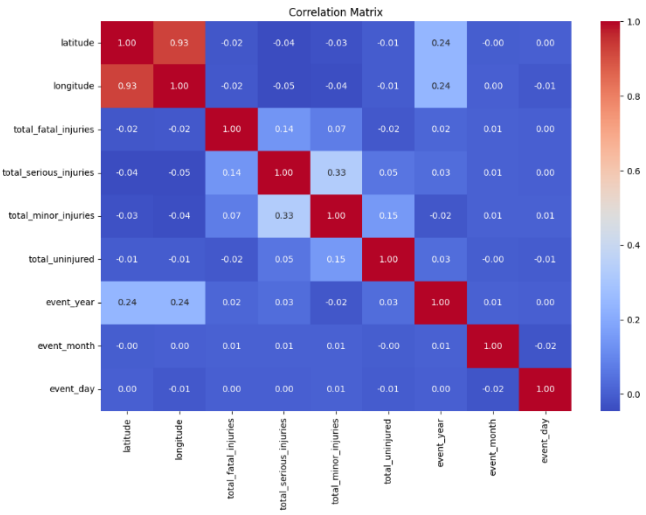

Conclusion

With these EDA operations, we identified significant insights which informed us about our data processing decisions, feature selections, and potentially feature engineering for downstream modeling. These operations served to make the structure of the dataset and its underlying relationships clear but also showed areas for further investigation and concentration in the safety improvement process within the aviation industry. The results from this EDA provide a basis for further, more intensive statistical analysis and predictive modeling with the view to strive for even safer aviation and more effective regulation.

THIS IS THE HEADING FOR A TABLE

[a] This is a table footnote.

You can cite your references in text by including the corresponding number, in square brackets [1]. If you need to cite a specific part of the source, you can include a page number [2, p. 13] or range [3, pp. 41–56].

ACKNOWLEDGMENTS

"Acknowledgment(s)" is spelled without an "e" after the "g" in American English.

As you can see, the formatting ensures that the text ends in two equal-sized columns rather than only displaying one column on the last page.

This template was adapted from those provided by the IEEE on their own website.

REFERENCES

[1] https://www.kaggle.com/datasets/khsamaha/aviation-accident-database-synopses?select=AviationData.csv